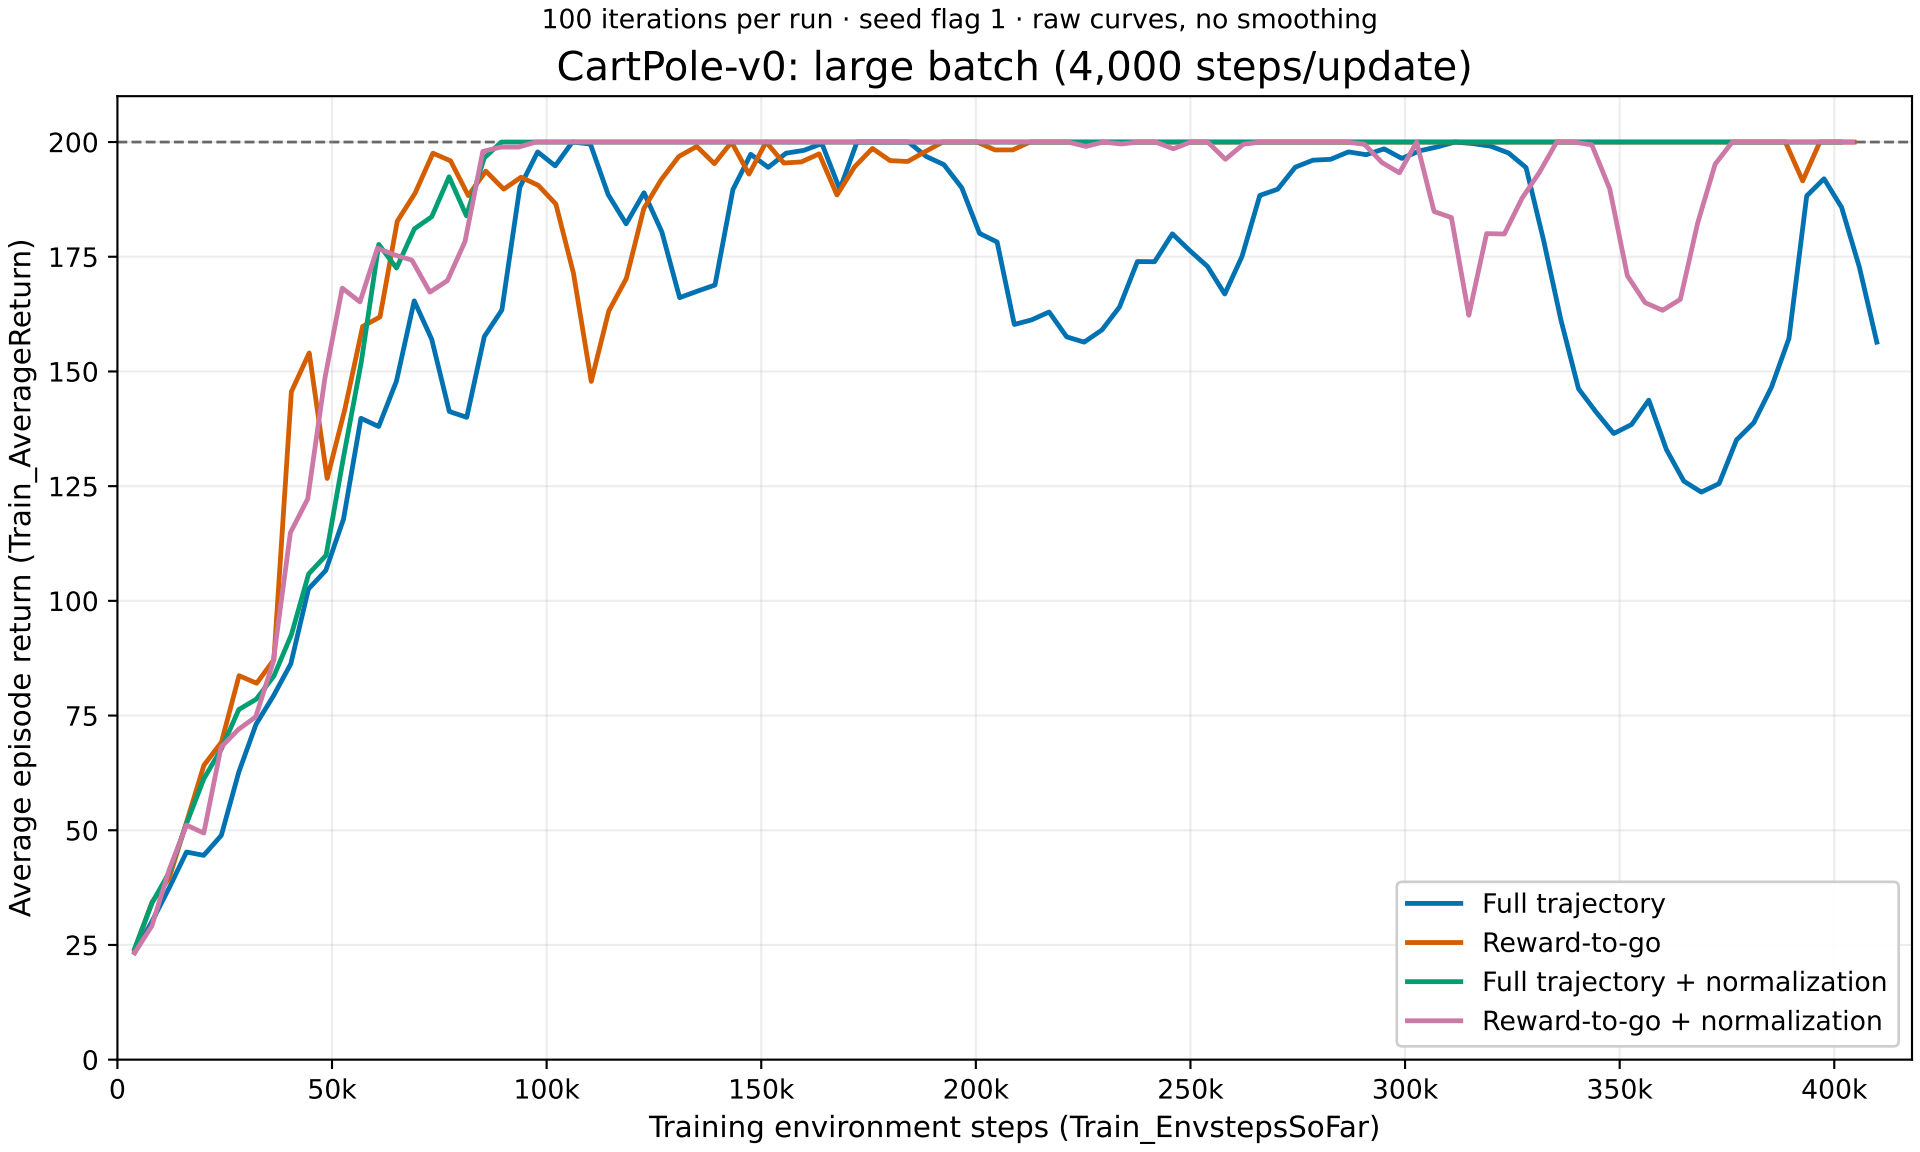
100 iterations per run · seed flag 1 · raw curves, no smoothing
CartPole-v0: large batch (4,000 steps/update)
Average episode return (Train_AverageReturn)
Training environment steps (Train_EnvstepsSoFar)
0
25
50
75
100
125
150
175
200
0
50k
100k
150k
200k
250k
300k
350k
400k
Full trajectory
Reward-to-go
Full trajectory + normalization
Reward-to-go + normalization